

CSC3121

Distributed Systems

Introduction and an Overview

Dr Paul Ezhilchelvan

Aim

- To expose the challenges in building distributed applications
- To understand the **concepts** behind the solutions
 - These solutions, on being implemented, form a part of **middleware** in distributed systems
 - SOAP, distributed transactions, etc.,
- Vendors implement these solutions
 - For various OS, process architecture
 - in various ways, in improved versions
 - Improvement requires conceptual understanding

Developing Distributed Apps

- What is a distributed system?
 - Autonomous Computers connected through network communication
 - LAN, Internet, Wireless...
- Developing a distributed application
 - Involves enabling a program in one computer to access 'resources' on other computers
- Middleware supports resource access in a **transparent** manner
 - *As though* the resources reside in the local computer
 - Creates a useful *illusion* to the developer

22 September 2020

3

Being Transparent

- Hide resource distribution from developers and their applications
- Location transparency
 - An application need not be aware of resource location
 - Middleware sorts it out
- Migration transparency
 - Resource may have been moved, acquiring a new address
 - Middleware finds the new address in the next access
- Concurrent access -> sequential access
- Failure transparency

22 September 2020

4

Access Transparency

- An app in computer C1 wants to read a file *fred* in C2
 - App is called the user or *client*
- If *fred* is local to C1, client can access it any time
- So, a program in C2 must be running 24/7
 - Dealing with access requests for *fred* from clients
 - It is called the *server*
- Say, client wants to
 - read *nbytes*
 - starting from position *pos*
 - Into **local** memory starting from address *buf*
- Client should do no more than what it would do, had *fred* been local

22 September 2020

5

Access transparency in programming

- Had *fred* been local, the client would execute:


```
read(buf, fred, nbytes, pos)
```
- Middleware in client machine
 - Is **informed** `read(.., fred, ..)` invocation cannot find *fred*
 - Creates a thread to deal with `read(.., fred, ..)`
 - The thread
 - locates *fred*
 - Turns the read instruction into a message
 - Sends the message to server machine
 - Puts the server response in *buf*
- Client programming is made easy by middleware
 - Remote procedure call (RPC)
 - Remote method invocation (RMI)
 - Simple Object Access protocol (SOAP)
 - Smart contracts in Hyperledger Blockchain systems

22 September 2020

6

Transaction: Money Transfer

- Alice has a bank account **A**
- Bank's computer is **C1**
- Alice instructs **C1** to transfer £250 to Bob's account **B**
- Bob's bank's computer is **C2**
- **C1** executes this program:
 - $A = A - 250$ (A is 'local')
 - $B = B + 250$ (B is **remote**, so use SOAP/RPC)
 - Involves C1 sending a SOAP or RPC message to C2

22 September 2020

7

Uncertainty Due to Crashes

- After sending RPC message but before receiving server response,
 - C1 detects that C2 had crashed
- Did C2 crash before or after doing $B = B + 250$?
- C1 cannot know this.. until C2 recovers
- Detecting this, and repairing any damage
 - are not Alice's (user) problems
 - Application deals with them (e.g., using a **Transaction Engine**)
- Had A and B been objects in a single computer, recovery is not a problem
 - **Atomic** write
 - Distribution transaction must be atomic

22 September 2020

8

Topics, Coursework & Pre-Requisite

- Topics covered:
 1. RPC
 2. Time, Clocks and Events in a Distributed System
 3. Atomic Transactions
 4. *Lock-Free Transactions??*
 5. *Computing a Global picture???*
- Module does not involve programming
- But has an important **pre-requisite**
 - Take notes / Ask questions
- I take delight in answering questions – an interesting part of my job!
- Assessment
 - **Two** exercises (summative)
 - Another design exercise in total order (formative)
 - very challenging exercises that test your understanding
 - Listen to the lecture videos and ask questions

22 September 2020

9

Something to think about..

- Computer C1's clock reads
11:31:53 when GMT = 11:30:00
- Computer C2's clock reads
11:28:16 also when GMT = 11:30:00
- C1 and C2 are Internet connected
- Alice proposes the following to synchronize both clocks:
 - C1 sends its clock reading to C2;
 - C2 adjusts its clock to the reading it received;
- Under what condition(s), clocks of C1 and C2 are **guaranteed** to show the **same** value at **the same GMT**?
 - Note: Clocks reading 12:00:00 when GMT = 12:10:00 meets the condition of showing the same reading at the same GMT

22 September 2020

10

References for Additional Reading

Chapters in Books:

- Andrew Tanenbaum and Marteen van Steen. *Distributed Systems: Principles and Paradigms*, [Second Edition](#), Pearson/Prentice Hall (Chapter 4) (https://www.dropbox.com/s/0s94rmfo8op40vp/Distributed%20Systems_TanenbaumBook.pdf?dl=0)
- Marteen van Steen and Andrew Tanenbaum. *Distributed Systems*. Third Edition. (Chapter 4) [Preliminary Version](#) (<https://komputasi.files.wordpress.com/2018/03/mvsteen-distributed-systems-3rd-preliminary-version-3-01pre-2017-170215.pdf>)
- George Coulouris, Jean Dillimore, Tim Kindberg and Gordon Blair. *Distributed Systems: Concepts and Design*. Fifth Edition, Addison Wesley. (Chapters 14-17) [Preliminary Version](#). (<http://bedford-computing.co.uk/learning/wp-content/uploads/2016/03/george-coulouris-distributed-systems-concepts-and-design-5th-edition.pdf>)

Technical papers:

- *Time, clocks, and the ordering of events in a distributed system*, L Lamport. Communications of the ACM. (<ftp://129.241.210.18/disk3/pub/vms/freevms/doc/lamport.pdf>)
- *Distributed Snapshots: Determining Global States of Distributed Systems*, M Chandy and L Lamport, ACM Trans. On Computer Systems.